

Grounded Document Q&A: should you fund it?

For: Northwind Cloud, Support & Operations · **Pattern:** RAG with citation grounding
Recommendation: fund a two-week pilot, with one condition (below).

The problem

Your support team answers the same questions out of the same documents all day. Refund windows, SLA credits, retention periods, SSO setup. They still answer them inconsistently. An assistant that reads those documents is the obvious fix.

The reason it isn't already done is the risk. An AI assistant that is wrong is usually wrong confidently, and nobody downstream can tell which answers to trust. One wrong refund figure quoted to a customer costs more than the tool saves in a month, and it ends the team's willingness to use it.

Citation grounding doesn't buy you better answers. It buys you the ability to tell a good answer from a bad one, before the customer sees it.

What “production-ready” means here

A demo retrieves a passage, writes a fluent answer and attaches a source. It takes a day and looks finished. Four things separate that from something you can put in front of customers:

	Demo	Production
Citations	attached	verified , so a claim whose source doesn't support it gets removed
Not knowing	answers anyway	refuses and says why ; “not in our docs” is a correct answer
Accuracy	unknown	measured against questions with agreed right answers
Documents	one clean folder	your real ones, contradictions and stale copies included

The naive version fails silently. With no way to say “I don't know” it always produces something, and you find out it was wrong when a customer does.

Three things to weigh before committing budget

- Your documents set the ceiling, and they're your responsibility.** Where two policies contradict, the system will faithfully cite the wrong one, and no amount of engineering fixes that. You'll need a named owner for the document set; not having one is the commonest reason these projects underdeliver.
- A citation proves provenance, not correctness.** Our verification catches the expensive error, where a claim states a figure its source doesn't contain: “14 days” against “30 days”. It won't catch an answer that paraphrases its source while inverting the meaning. Closing that gap roughly doubles the cost per answer, so we'd rather decide it after the pilot, on real data.
- Launch isn't the finish.** Documents change and accuracy decays quietly. Budget for ongoing evaluation and re-indexing, or in six months you'll have a tool people have stopped trusting without anyone able to say when.

Two-week scope

Week 1, measurement before engineering. We ingest your real documents and tune retrieval on them, and build a 50-question evaluation set with agreed correct answers. That set de-risks everything after it, because it's how both of us find out whether this works.

Week 2, calibration and pilot. We tune citation verification and the refusal threshold against that set. Whether it should answer more or be wrong less is your call, and we'll bring the numbers. Then a pilot with five to ten agents, and handover.

In scope: one document set, one internal interface, English, read-only. **Out of scope:** other languages, acting on a customer's behalf, live document sync, permission-scoped answers. Each is a sensible phase two, and adding one now puts the pilot at risk.

From you: a document-set owner, five to ten pilot users, and three hours of a subject expert's time to agree the evaluation answers.

The condition

Fund the pilot, but treat the week-one evaluation set as a go/no-go gate. If accuracy on your own documents comes back low, the documents are the problem rather than the model, and fixing them is where the next money should go. We'd rather tell you that in week two than in month six.