

Grounded Document Q&A — should you fund it?

For: Northwind Cloud, Support & Operations · **Pattern:** RAG with citation grounding
Recommendation: fund a two-week pilot, with one condition (below).

The problem, plainly

Your support team answers the same questions out of the same documents all day — refund windows, SLA credits, retention periods, SSO setup — and still answers them inconsistently. An assistant that reads those documents is the obvious fix.

The reason it isn't already done is the risk: **an AI assistant that is wrong is wrong confidently, and nobody downstream can tell which answers to trust.** One wrong refund figure quoted to a customer costs more than the tool saves in a month, and it ends the team's willingness to use it at all.

Citation grounding does not buy better answers. It buys **the ability to tell a good answer from a bad one** — before the customer sees it.

What “production-ready” means here, and what a demo hides

A demo retrieves a passage, writes a fluent answer and attaches a source. It takes a day and it looks finished. Four things separate it from something customer-facing:

	Demo	Production
Citations	attached	verified — a claim whose source doesn't support it is removed
Not knowing	answers anyway	refuses, and says why ; “not in our docs” is a correct answer
Accuracy	unknown	measured against questions with agreed right answers
Documents	one clean folder	your real ones, contradictions and stale copies included

The naive version **fails silently**: with no way to say “I don't know”, it always produces something. You find out it was wrong when a customer does.

Three things to weigh before committing budget

- Your documents set the ceiling, and they are your responsibility.** Where two policies contradict, the system faithfully cites the wrong one; no engineering fixes that. You need a named owner for the document set — its absence is the commonest reason these projects underdeliver.
- A citation proves provenance, not correctness.** Verification catches the common, expensive error — a figure the source doesn't contain, “14 days” against “30 days” — but not an answer that paraphrases its source while inverting the meaning. Closing that gap roughly doubles the per-answer cost; decide after the pilot, on real data.
- Launch is not the finish.** Documents change and accuracy decays quietly. Budget for ongoing evaluation and re-indexing, or in six months you have a tool people have quietly stopped trusting.

Two-week scope

Week 1 — measurement before engineering. Ingest and tune retrieval on your real documents, and build a **50-question evaluation set** with agreed correct answers. This is the deliverable that de-risks the rest: it is how both of us find out whether this works.

Week 2 — calibration and pilot. Tune citation verification and the refusal threshold against that set. The trade-off between “answers more” and “is wrong less” is your decision; we bring the numbers. Then a pilot with 5–10 agents, and handover.

In scope: one document set, one internal interface, English, read-only. **Explicitly out:** other languages, taking actions for a customer, live document sync, permission-scoped answers. Each is a sensible phase two; adding one now puts the pilot at risk.

From you: a document-set owner, 5–10 pilot users, and about three hours of a subject expert's time to agree the evaluation answers.

The condition

Fund the pilot, and treat **the week-one evaluation set as the go/no-go gate**. If accuracy on your own documents comes in low, the problem is the documents, not the model — and the right next spend is fixing them, not more engineering. We would rather tell you that in week two than in month six.